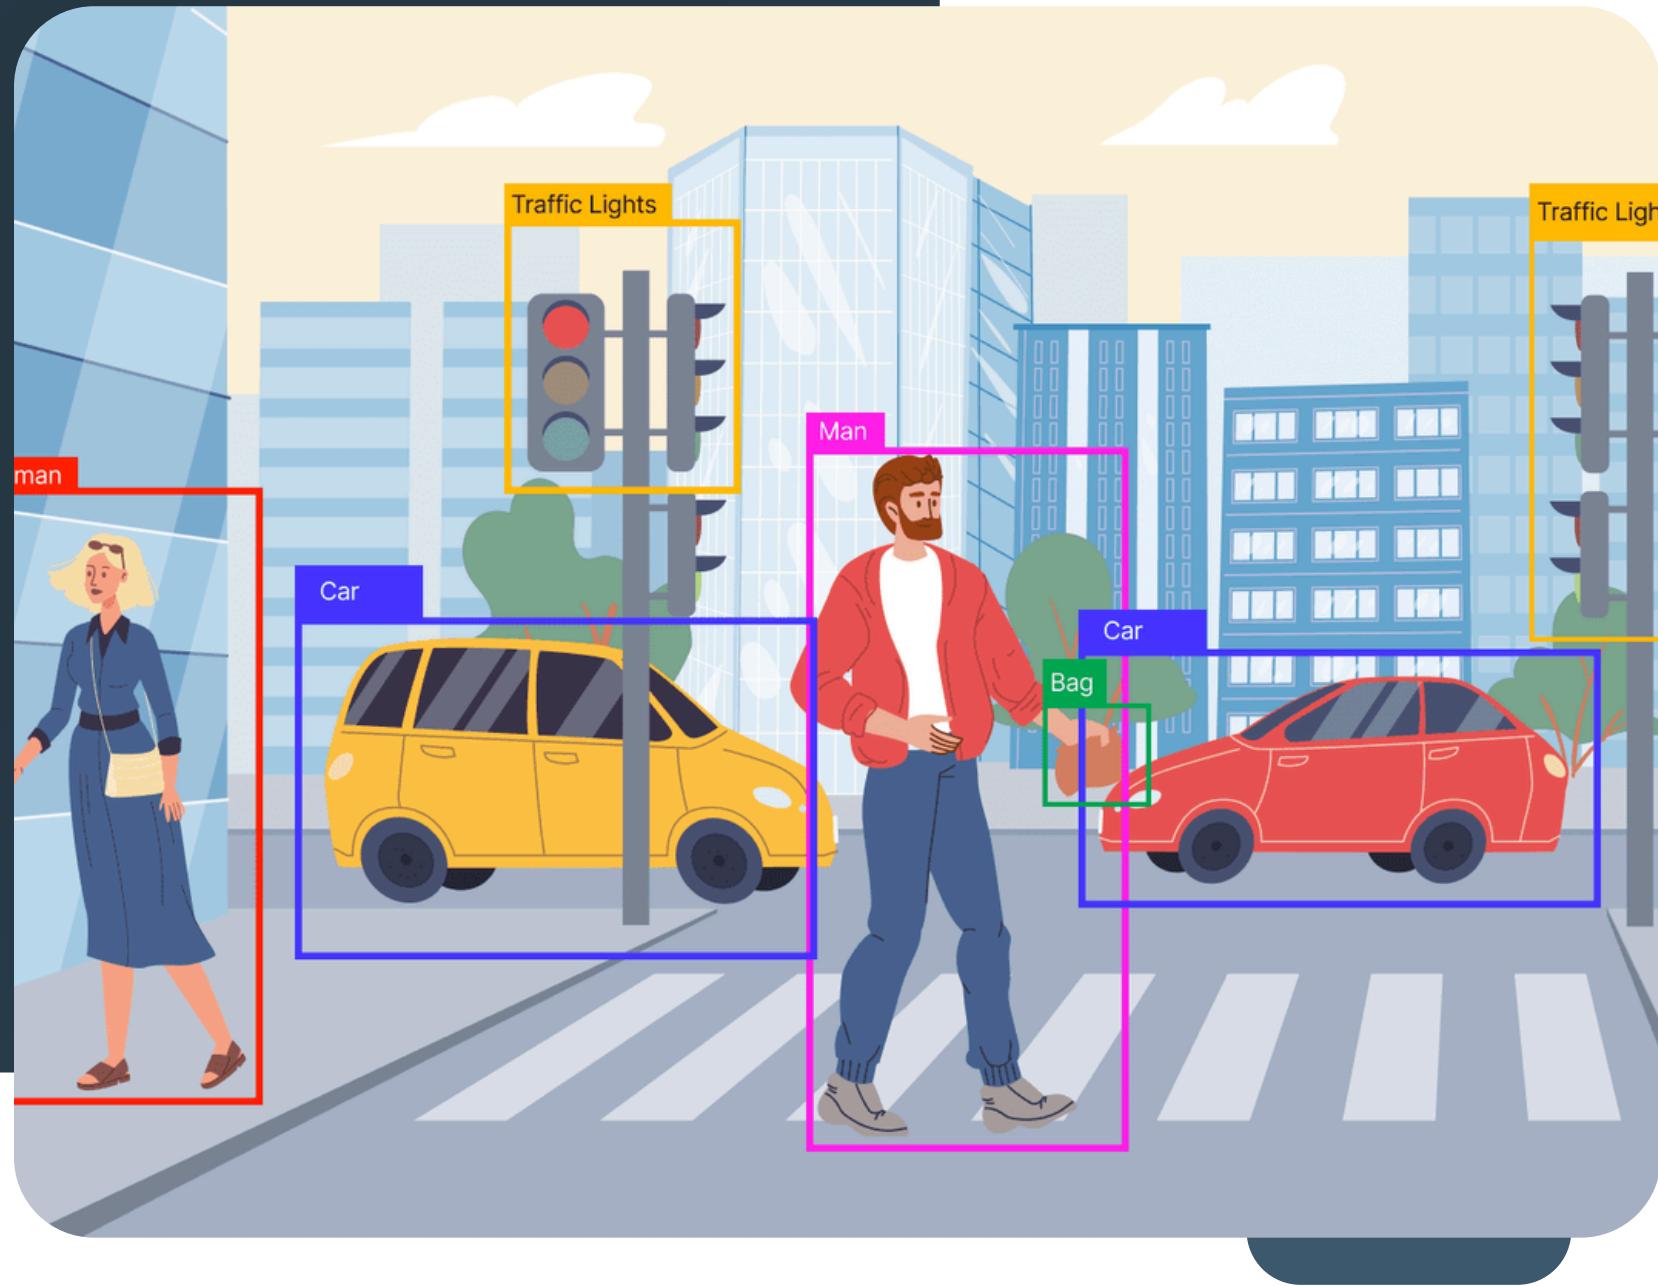


NEURAL NETWORK MODELS FOR OBJECT RECOGNITION

CIFAR-10

● KANAN AKPAROV



WHY OBJECT RECOGNITION MATTERS

A core part of modern AI, powering:

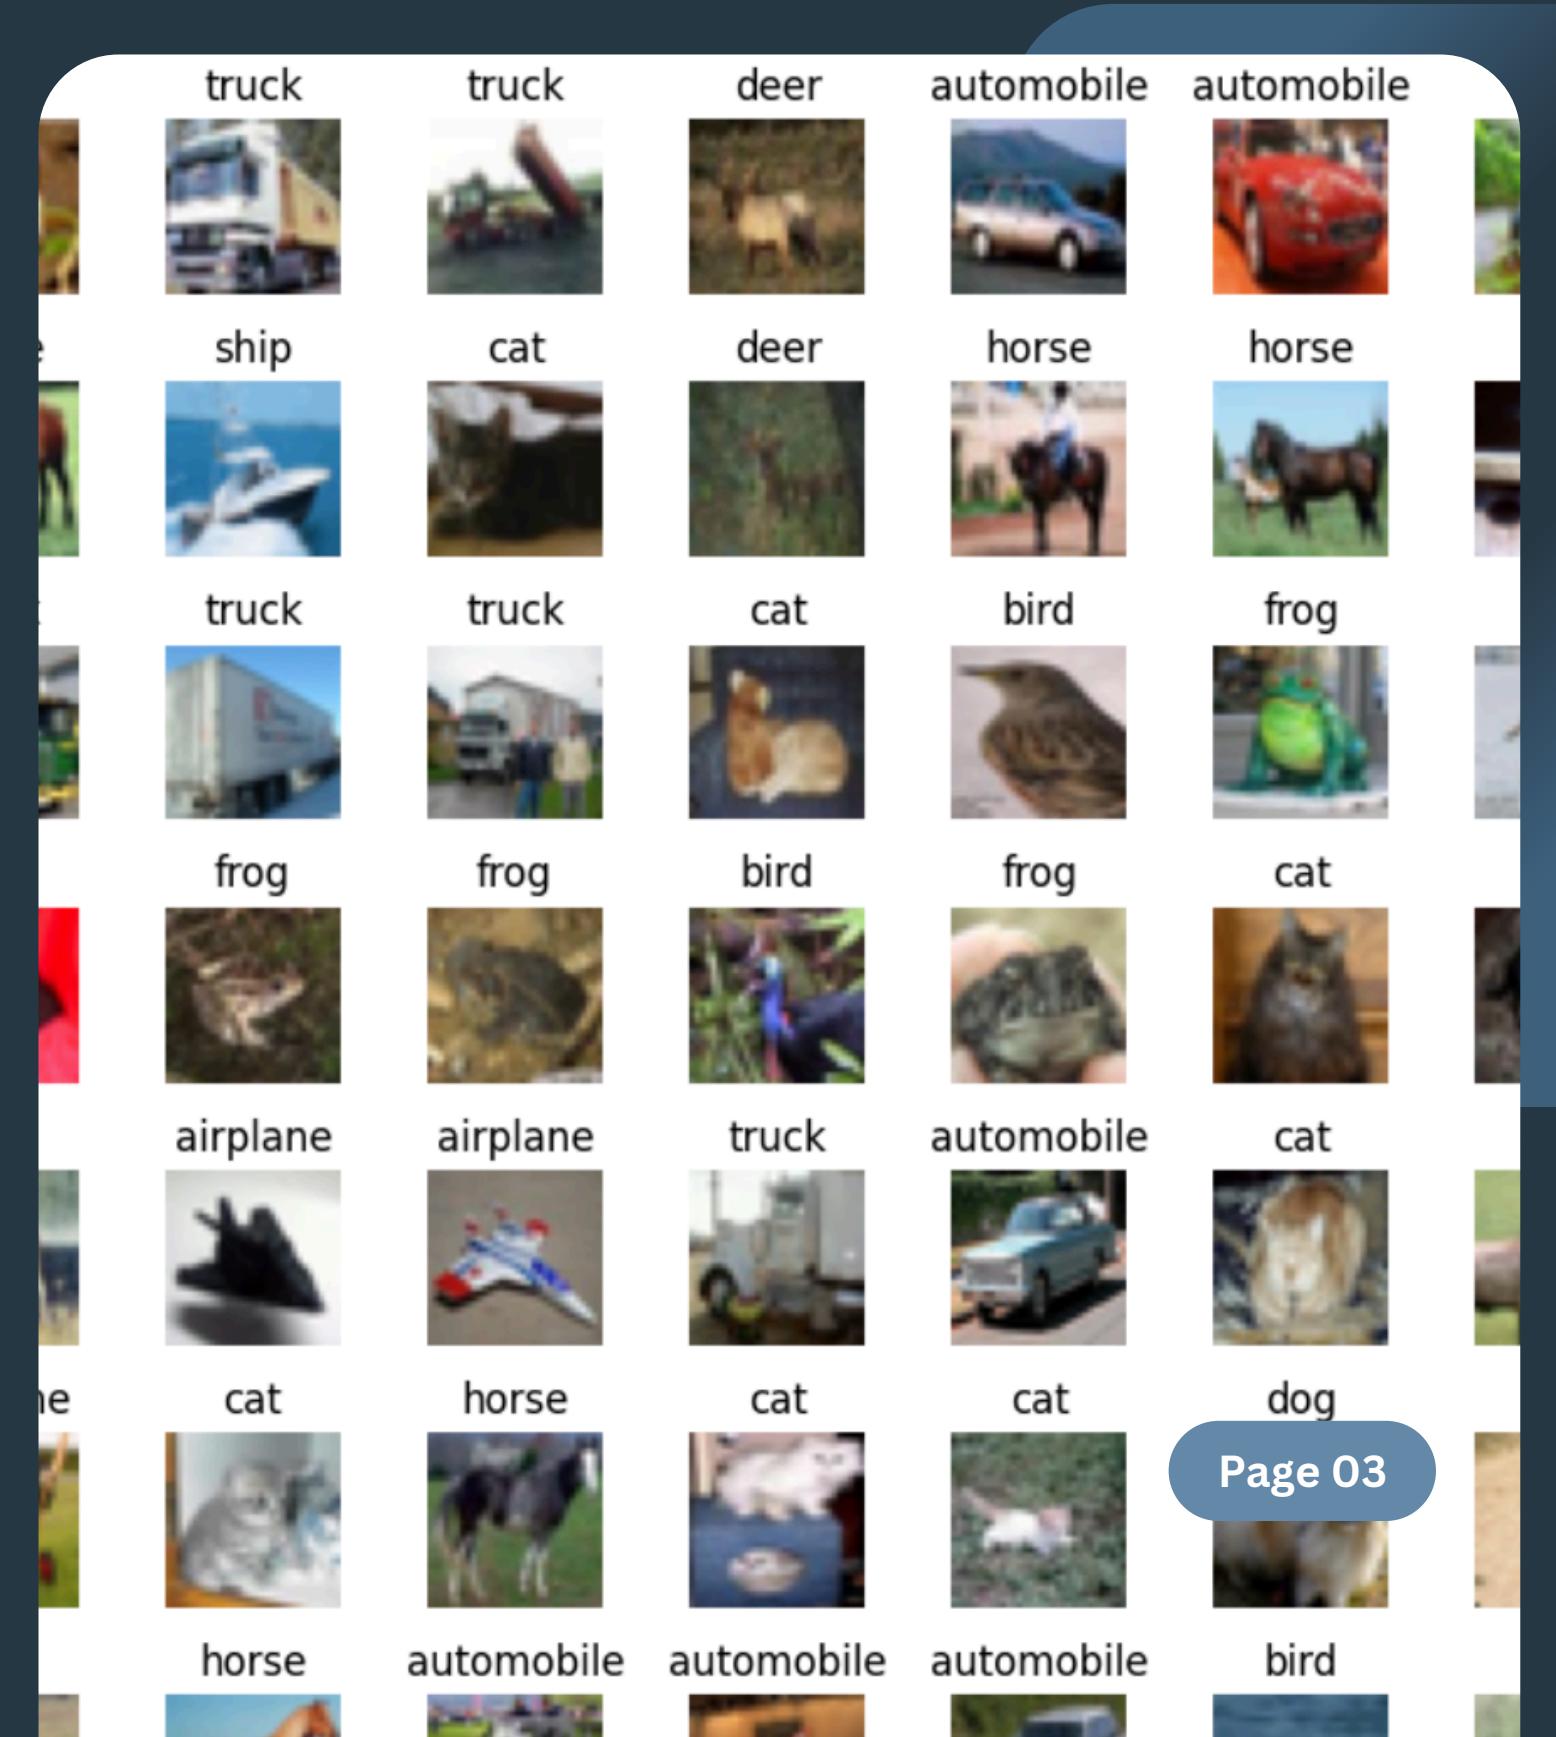
- Smart vehicles
- Facial recognition
- Quality control
- Healthcare imaging
- Retail automation

DATASET OVERVIEW

The dataset contains 60,000 color images with a resolution of 32×32 pixels, split into:

01.

- 50,000 for training
- 10,000 for testing
- 5,000 reserved for validation



```
import numpy as np
import matplotlib.pyplot as plt
from tensorflow.keras.datasets import cifar10
from tensorflow.keras.utils import to_categorical

class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer',
               'dog', 'frog', 'horse', 'ship', 'truck']

(X_train, y_train), (X_test, y_test) = cifar10.load_data()
X_train = X_train.astype('float32') / 255.0
X_test = X_test.astype('float32') / 255.0

# One-hot encode labels
y_train_cat = to_categorical(y_train, 10)
y_test_cat = to_categorical(y_test, 10)

print("Training data shape:", X_train.shape, y_train.shape)
print("Test data shape:", X_test.shape, y_test.shape)
print("Number of classes:", len(class_names))

# Display
def plot_sample_images(X, y, class_names):
    plt.figure(figsize=(8, 8))
    for i in range(49):
        plt.subplot(7, 7, i+1)
        plt.imshow(X[i])
        plt.title(class_names[int(y[i])])
        plt.axis('off')
    plt.tight_layout()
    plt.show()
plot_sample_images(X_train, y_train, class_names)
```

DATASET METADATA & SPLITTING

The original 50,000 images are split to evaluate model generalization on unseen data.

Each image has 3,072 features, providing comprehensive input for learning.

IMPORTANCE OF VALIDATION

- Offers feedback during model development
- Helps tune parameters like learning rate and dropout rate
- Detects overfitting
- Provides an unbiased assessment of model performance

```
# Perform stratified sampling to split training into training and validation sets
X_train, X_val, y_train, y_val = train_test_split(
    X_train, y_train,
    test_size=0.2,
    random_state=42,
    stratify=y_train # Stratify by labels to preserve class distribution
)

# Build and compile the CNN model
cnn = models.Sequential([
    layers.Conv2D(filters=32, kernel_size=(3, 3), activation='relu', input_shape=(32, 32, 3)),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(filters=64, kernel_size=(3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(filters=128, kernel_size=(3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Flatten(),
    layers.Dense(60, activation='relu'),
    layers.Dense(40, activation='relu'),
    layers.Dense(10, activation='softmax')
])

cnn.compile(optimizer='adam',
            loss='sparse_categorical_crossentropy',
            metrics=['accuracy'])

# Train the model and save the history
history = cnn.fit(X_train, y_train, epochs=10, validation_data=(X_val, y_val))

# Function to plot training and validation loss and accuracy
def plot_training_history(history):
    # Retrieve metrics from the history object
    acc = history.history['accuracy']
    val_acc = history.history['val_accuracy']
    loss = history.history['loss']
    val_loss = history.history['val_loss']
    epochs = range(1, len(acc) + 1)

    # Plot accuracy
    plt.figure(figsize=(12, 5))
    plt.subplot(1, 2, 1)
    plt.plot(epochs, acc, label='Training Accuracy')
    plt.plot(epochs, val_acc, label='Validation Accuracy')
    plt.title('Training and Validation Accuracy')
    plt.xlabel('Epochs')
    plt.ylabel('Accuracy')
    plt.legend()

    # Plot loss
    plt.subplot(1, 2, 2)
    plt.plot(epochs, loss, label='Training Loss')
```

CNN MODEL ARCHITECTURE OVERVIEW & KEY MODEL COMPONENTS

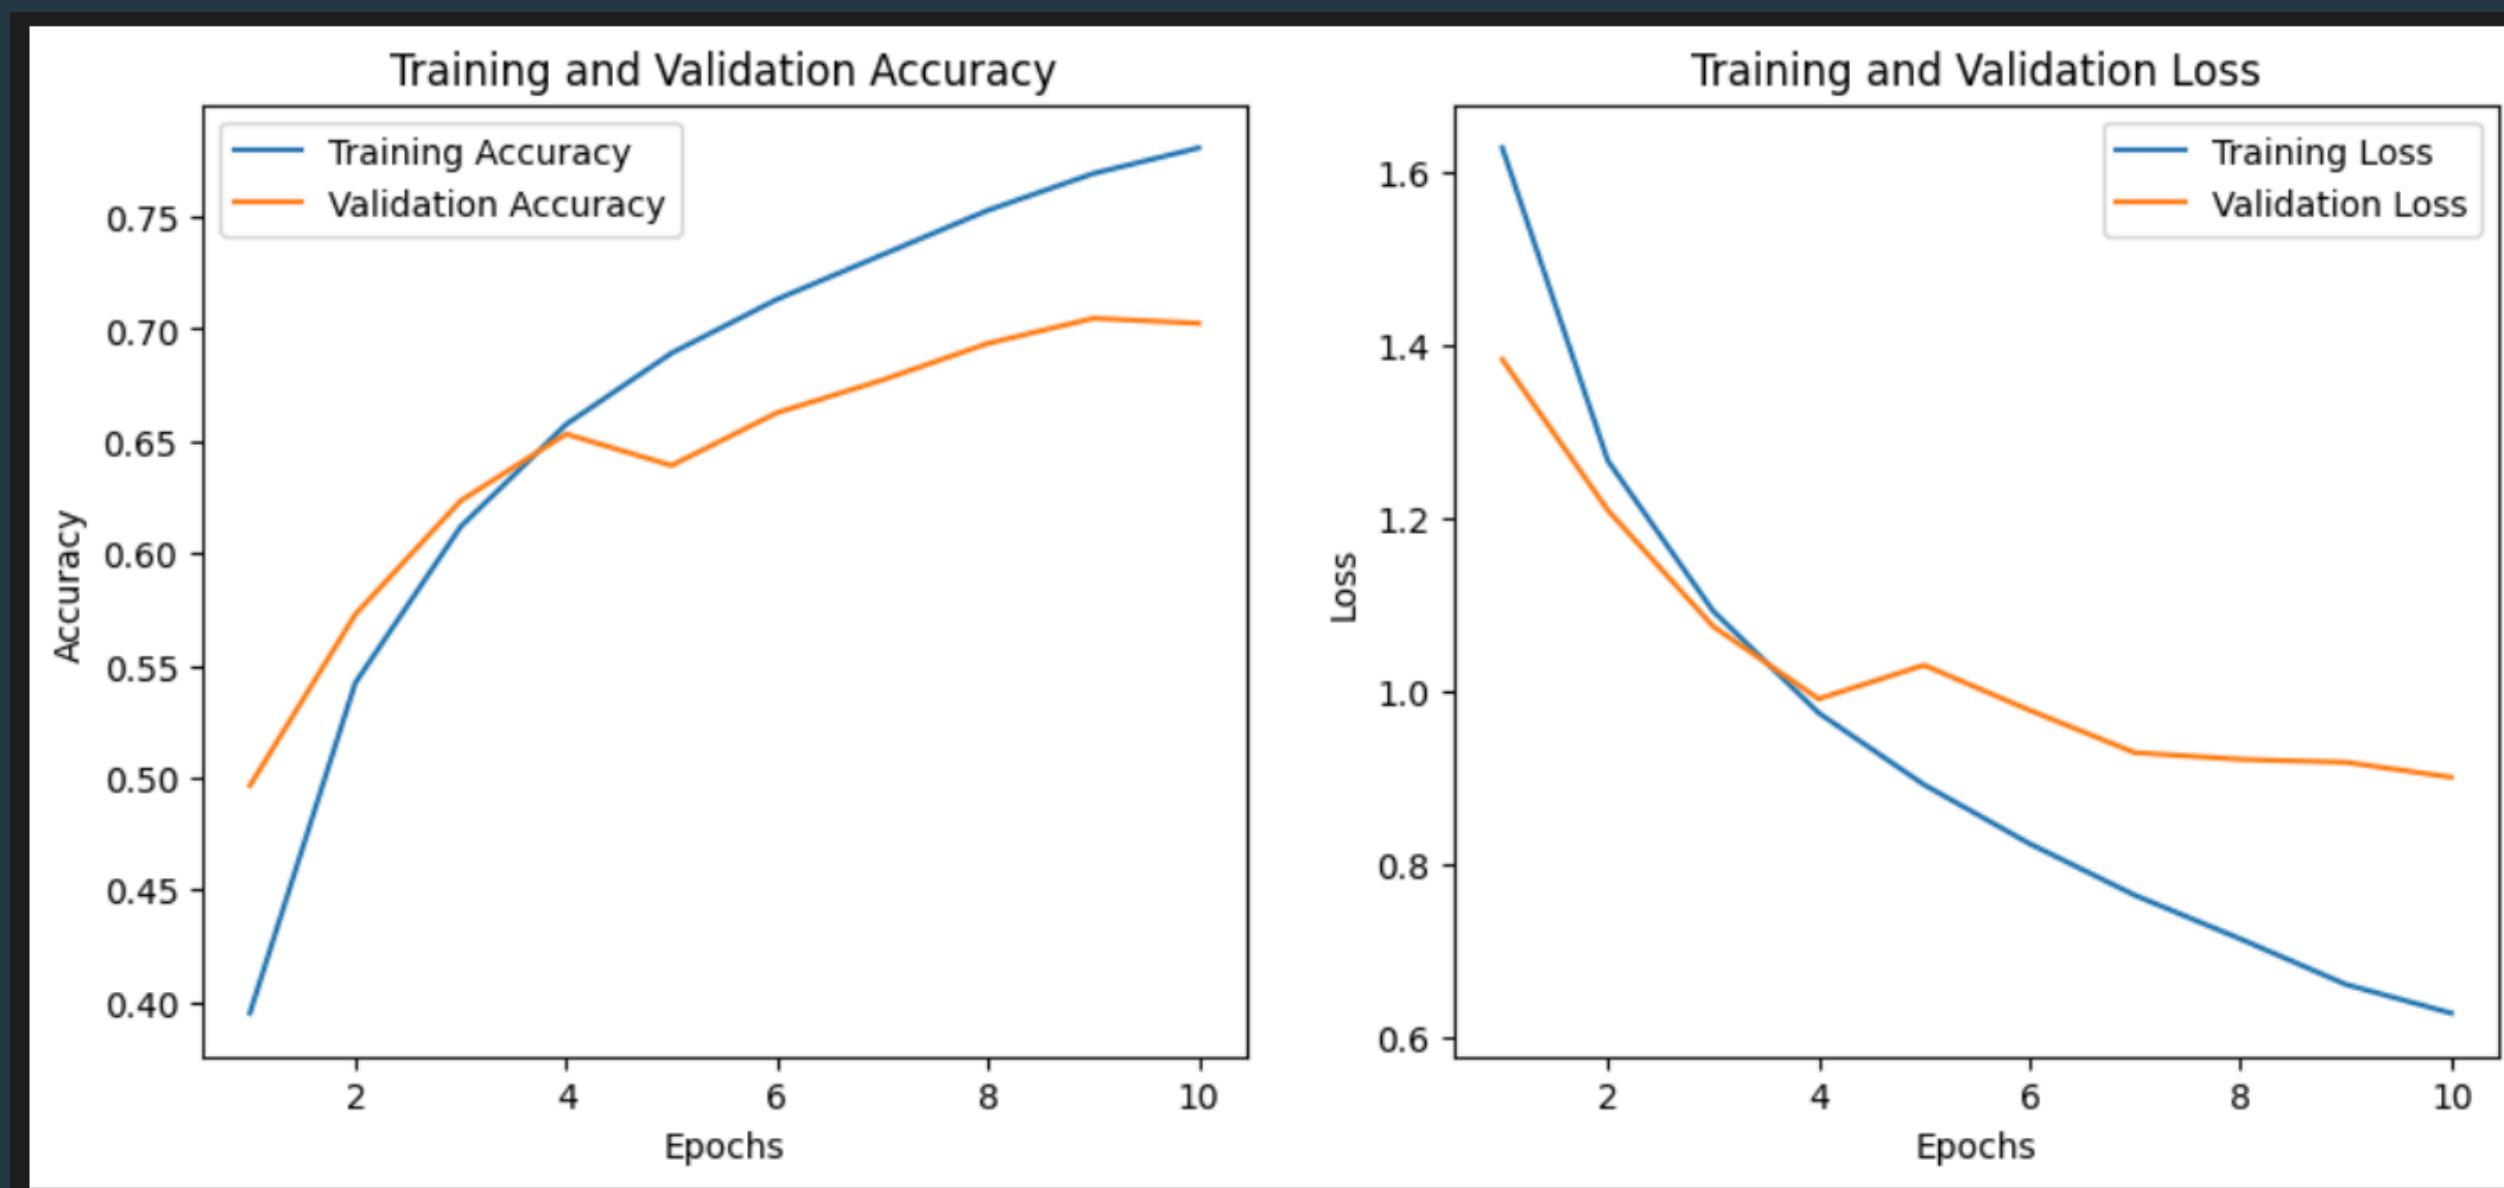
```
X_train, X_val, y_train, y_val = train_test_split(  
    X_train, y_train,  
    test_size=0.2,  
    random_state=42,  
    stratify=y_train # Stratify by labels to preserve class distribution  
)  
  
# Build and compile the CNN model  
cnn = models.Sequential([  
    layers.Conv2D(filters=32, kernel_size=(3, 3), activation='relu', input_shape=(32, 32, 3)),  
    layers.MaxPooling2D((2, 2)),  
    layers.Conv2D(filters=64, kernel_size=(3, 3), activation='relu'),  
    layers.MaxPooling2D((2, 2)),  
    layers.Conv2D(filters=128, kernel_size=(3, 3), activation='relu'),  
    layers.MaxPooling2D((2, 2)),  
    layers.Flatten(),  
    layers.Dense(60, activation='relu'),  
    layers.Dense(40, activation='relu'),  
    layers.Dense(10, activation='softmax')  
)  
  
cnn.compile(optimizer='adam',  
            loss='sparse_categorical_crossentropy',  
            metrics=['accuracy'])  
  
# Train the model and save the history  
history = cnn.fit(X_train, y_train, epochs=10, validation_data=(X_val, y_val))  
  
# Function to plot training and validation loss and accuracy  
def plot_training_history(history):  
    # Retrieve metrics from the history object  
    acc = history.history['accuracy']  
    val_acc = history.history['val_accuracy']  
    loss = history.history['loss']  
    val_loss = history.history['val_loss']  
    epochs = range(1, len(acc) + 1)  
  
    # Plot accuracy  
    plt.figure(figsize=(12, 5))  
    plt.subplot(1, 2, 1)  
    plt.plot(epochs, acc, label='Training Accuracy')  
    plt.plot(epochs, val_acc, label='Validation Accuracy')  
    plt.title('Training and Validation Accuracy')  
    plt.xlabel('Epochs')  
    plt.ylabel('Accuracy')  
    plt.legend()  
  
    # Plot loss  
    plt.subplot(1, 2, 2)
```

- Input: 32×32×3 RGB images (CIFAR-10)
- 3 Convolutional layers (32, 64, 128 filters) – extract visual features
- MaxPooling after each Conv – reduces spatial size
- Flatten layer – converts feature maps to vector
- Dense layers (60, 40 neurons) – learn complex patterns
- Output layer (10 neurons, softmax) – class prediction
- Optimizer: Adam | Loss: Categorical Crossentropy | Metric: Accuracy
- Trained for 10 epochs with validation split (20%)
- Plotted accuracy & loss to monitor model performance

TRAINING PARAMETERS

- Optimizer: Adam
- Loss Function: Categorical Crossentropy
- Metric: Accuracy
- Epochs: 10
- Validation Split: 20% (using `train_test_split`)
- Batch Size: (default – likely 32)
- Input Shape: (32, 32, 3)
- Output Classes: 10

ACCURACY AND LOSS OVER EPOCHS



TEST SET EVALUATION

- Final Accuracy: 70%
- Model performed well but can be improved

13/313 ————— 2s 6ms/step

Test Loss: 0.9032906293869019

Test Accuracy: 0.698300040054321

Test Precision: 0.706622169683003

Test Recall: 0.6982999999999999

Test F1-score: 0.6965683529846298

Classification Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| airplane | 0.70 | 0.80 | 0.75 | 1000 |
| automobile | 0.81 | 0.84 | 0.83 | 1000 |
| bird | 0.71 | 0.49 | 0.58 | 1000 |
| cat | 0.47 | 0.58 | 0.52 | 1000 |
| deer | 0.71 | 0.55 | 0.62 | 1000 |
| dog | 0.63 | 0.57 | 0.60 | 1000 |
| frog | 0.81 | 0.73 | 0.77 | 1000 |
| horse | 0.62 | 0.84 | 0.72 | 1000 |
| ship | 0.79 | 0.82 | 0.80 | 1000 |
| truck | 0.80 | 0.76 | 0.78 | 1000 |
| accuracy | | | 0.70 | 10000 |
| macro avg | 0.71 | 0.70 | 0.70 | 10000 |
| weighted avg | 0.71 | 0.70 | 0.70 | 10000 |

True: bird, Pred: bird

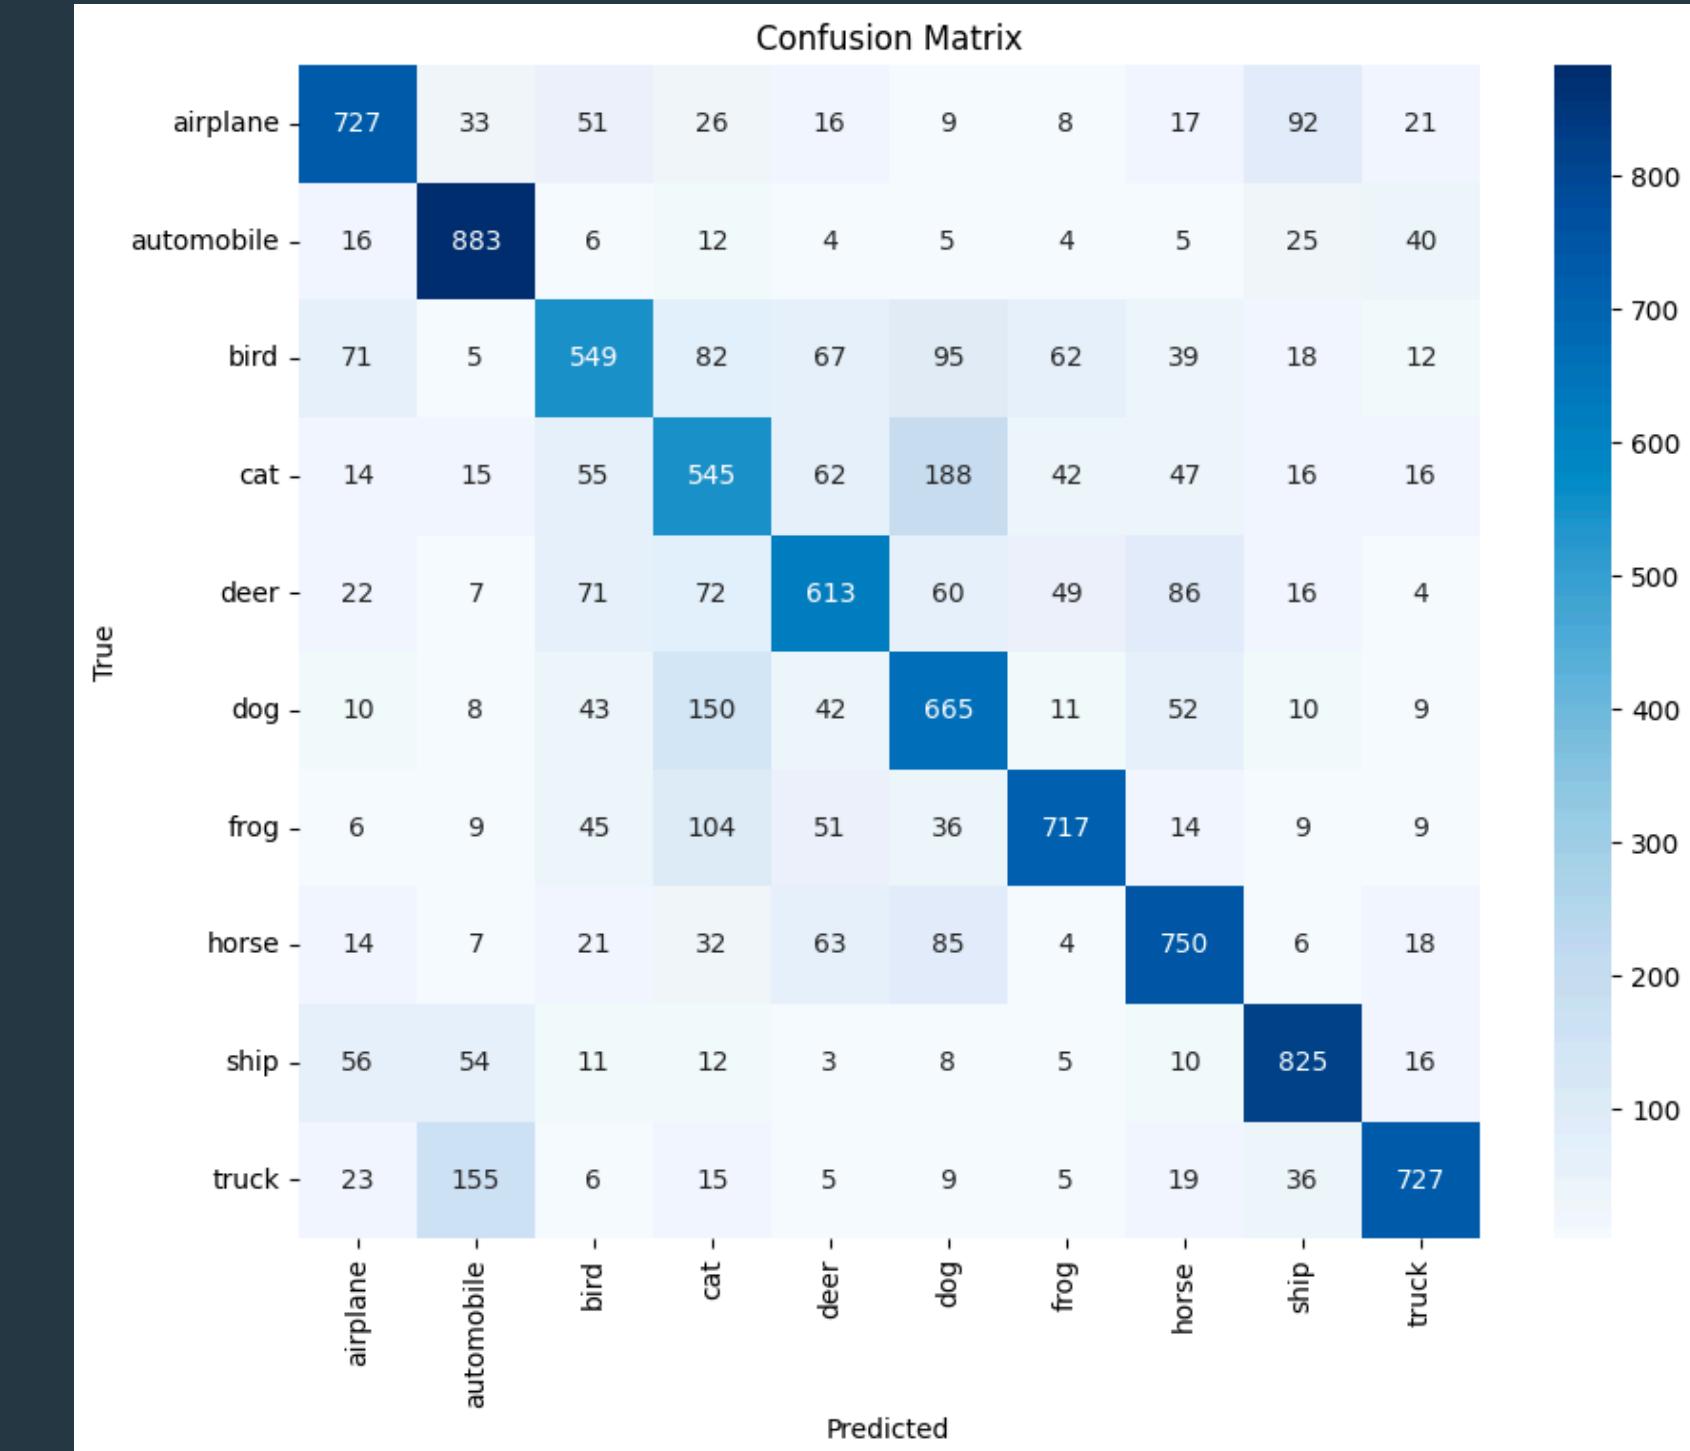


True: frog, Pred: cat



CONFUSION MATRIX & CLASS ANALYSIS

- Highest accuracy achieved on trucks, automobile, frogs, airplane
- Lowest accuracy observed on cats and birds
- Potential improvements: address underperforming classes



POSSIBLE IMPROVEMENTS

- Add more convolutional layers
- Apply Batch Normalization for stability
- Use Transfer Learning with a pre-trained model
(e.g., VGG16, ResNet)

```
from tensorflow.keras import models, layers, applications
import matplotlib.pyplot as plt

# Pre-trained VGG16 as feature extractor
base_model = applications.VGG16(weights='imagenet', include_top=False, input_shape=(32,32,3))
base_model.trainable = False

# Improved CNN model
cnn_improved = models.Sequential([
    base_model,
    layers.GlobalAveragePooling2D(),
    layers.Dense(128, activation='relu'),
    layers.BatchNormalization(),
    layers.Dense(10, activation='softmax')
])

# Compile the model
cnn_improved.compile(
    optimizer='adam',
    loss='categorical_crossentropy',
    metrics=['accuracy']
)

# Train the model
history_improved = cnn_improved.fit(
    X_train, y_train,
    epochs=10,
    validation_data=(X_val, y_val),
    batch_size=32
)
```

DESIGN DECISIONS & STRATEGY

- Built a simple yet effective CNN for CIFAR-10
- Prototype has enough depth to learn features while remaining interpretable
- Explored dropout, batch normalization, and increased filter counts to improve performance
- Consider future improvements: transfer learning, automated hyperparameter tuning, and advanced regularization techniques

CHALLENGES ENCOUNTERED

- Balanced training time versus model complexity
- Tuned batch size and learning rate for stability
- Observed occasional overfitting on some classes
- Some noisy or inconsistent predictions in underperforming categories
- Memory management

REAL WORLD APPLICATIONS

- Systems for recognizing and interpreting visual information
- Tools for diagnosing medical conditions
- Monitoring and managing agricultural environments
- Enhancing roadway safety
- Ensuring security at airports
- Starting with minor, focused tasks that scale into solving more complex, high-impact challenges

REFLECTIONS AND LEARNING OUTCOMES

- Connects conceptual understanding with hands-on application
- Highlights the critical role of testing, validation, and performance assessment
- Emphasizes that machine learning is an iterative process, not the pursuit of a single flawless model
-

REFERENCES

- Abadi, M. et al., 2016. TensorFlow: A system for large-scale machine learning. [online] Available at: <https://www.tensorflow.org> [Accessed 12 October 2025].
- Chollet, F. et al., 2015. Keras. [online] Available at: <https://keras.io> [Accessed 12 October 2025].
- Goodfellow, I., Bengio, Y. & Courville, A., 2016. Deep Learning. Cambridge, MA: MIT Press.
- Krizhevsky, A., Nair, V. & Hinton, G., 2014. The CIFAR-10 Dataset. [online] Available at: <https://www.cs.toronto.edu/~kriz/cifar.html> [Accessed 12 October 2025].
- LeCun, Y., Bottou, L., Bengio, Y. & Haffner, P., 1998. Gradient-based learning applied to document recognition. Proceedings of the IEEE, 86(11), pp.2278–2324.

THANK YOU
FOR YOUR ATTENTION

